



Airbnb in Southbank, Melbourne

Presented by

Garfield Consulting Service

Guorong Liu, Fangwen Wang

Duosi Zheng, Yuanzhe Zhuang

Metropolitan College | Boston University | Dec 10, 2019



Agenda

- 1 Data Preparation & Exploration**
- 2 Prediction: Multiple Regression Model**
- 3 Classification**
 - A. K-Nearest Neighbors
 - B. Naïve Bayes
 - C. Classification Tree
- 4 Clustering**
- 5 Conclusion & Suggestions**




1. Data Preparation & Exploration: Missing Values



Methodologies of Data Cleaning:

- I. dropping NA values, N/A values, and blank cells for most of the data models
- II. mean values : substitute the values of missing values (i.e. clustering model)
- III. Use `sum(is.na(df))` and `md.pattern()` to check null values
- IV. Dummy variables \leftarrow categorical values

accommodates review scores cleaning data with higher importance



(example for checking a random data frame)

```
# filter the only neighborhood Southbank
df1 <- filter(df, neighborhood=="Southbank")
dim(df1)
## [1] 1248 84
1248*84
## [1] 104832
#tell us the missing values,
this gives the number of null values:
sum(is.na(df1))
## [1] 5156
```

1. Data Preparation & Exploration: Summary Statistic & Visualizations



Who can tell any words?



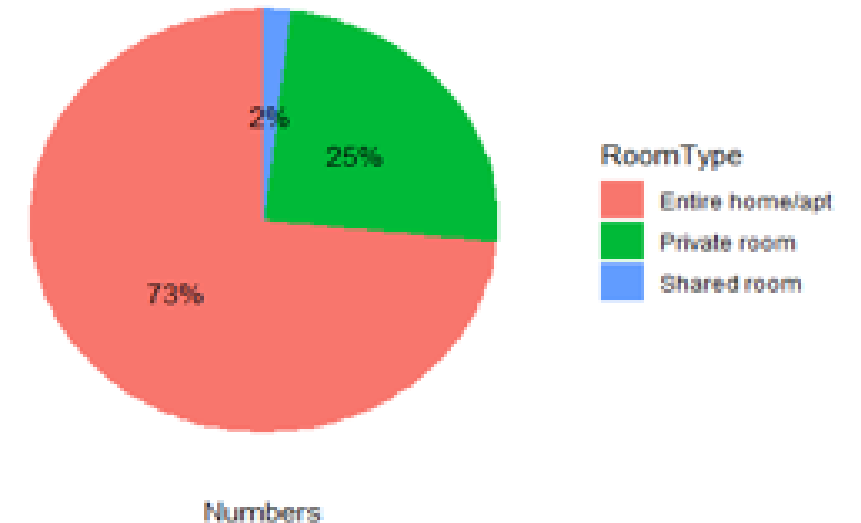
Word Cloud

Summaries of KPIs

[i.e. Review Scores Value (total:10)]

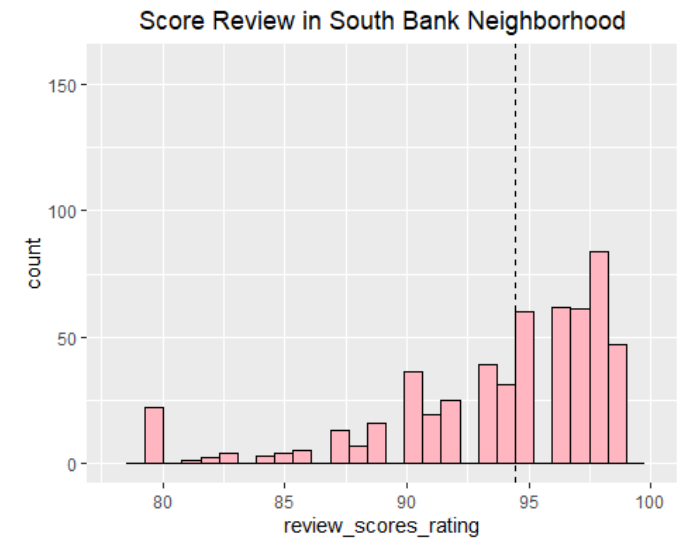



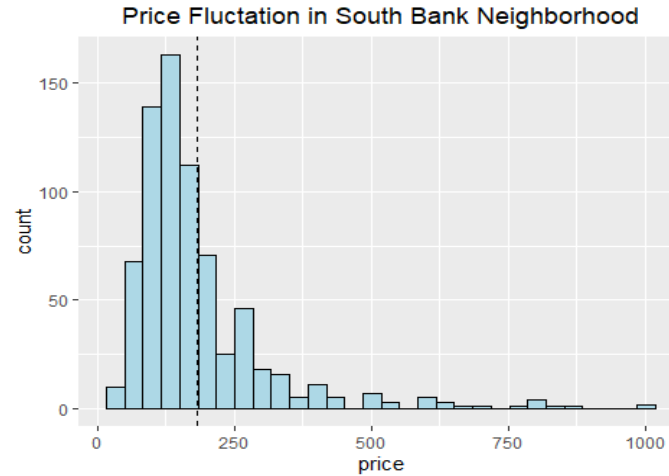
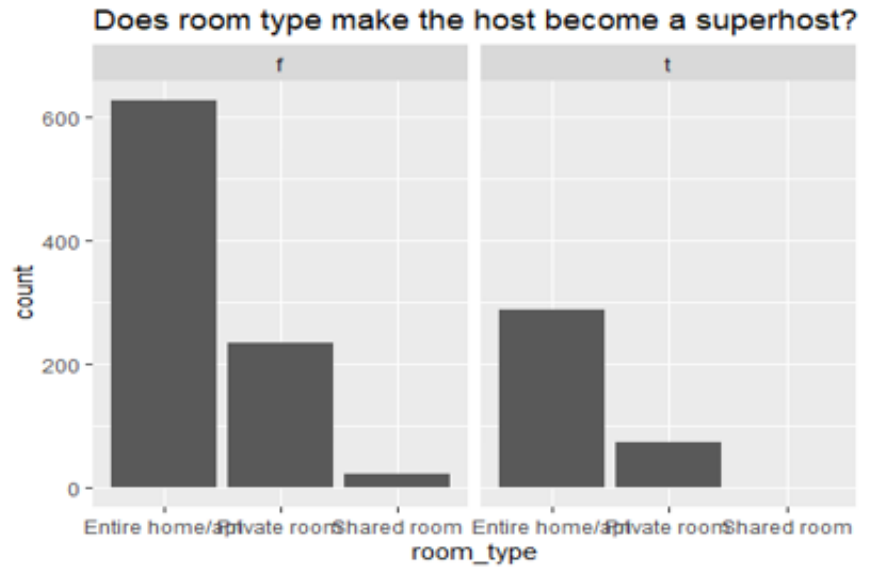
Mean	9.48
Standard Deviation	0.85
Median	10
Maximum	10
Minimum	2

Proportions of Properties



Most of the rooms are **entire home/apt.**

1. Data Preparation & Exploration: Summary Statistic & Visualizations



2. Prediction: Multiple Regression Model

```
> summary(model1)
```

```
Call:
lm(formula = price ~ ., data = train)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-185.85  -46.64  -13.40   26.39   910.49
```

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-29.09764	54.67879	-0.532	0.594880
accommodates	21.85964	3.91621	5.582	4.11e-08 ***
bathrooms	57.32668	11.19271	5.122	4.49e-07 ***
cleaning_fee	0.16204	0.12196	1.329	0.184619
host_is_superhost	-7.61686	10.32731	-0.738	0.461174
guests_included	9.70122	3.72047	2.608	0.009422 **
number_of_reviews	0.05668	0.11195	0.506	0.612888
review_scores_value	1.37148	5.60424	0.245	0.806783
reviews_per_month	-13.85994	3.14882	-4.402	1.34e-05 ***
security_deposit	0.02743	0.01488	1.843	0.065946 .
instant_bookable	5.49858	9.86502	0.557	0.577543
availability_30	1.95821	0.53451	3.664	0.000278 ***
calculated_host_listings_count	-0.79404	0.25108	-3.162	0.001670 **

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 98.19 on 452 degrees of freedom
Multiple R-squared:  0.4363,    Adjusted R-squared:  0.4213
F-statistic: 29.15 on 12 and 452 DF,  p-value: < 2.2e-16
```

```
> summary(model2)
```

```
Call:
lm(formula = price ~ accommodates + bathrooms + guests_included +
    reviews_per_month + security_deposit + availability_30 +
    calculated_host_listings_count, data = train)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-190.93  -47.03  -14.38   27.54   900.42
```

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-14.17781	15.54291	-0.912	0.362159
accommodates	24.10600	3.58895	6.717	5.54e-11 ***
bathrooms	57.58753	11.04951	5.212	2.84e-07 ***
guests_included	9.35797	3.64854	2.565	0.010640 *
reviews_per_month	-13.03112	2.44978	-5.319	1.63e-07 ***
security_deposit	0.02877	0.01448	1.987	0.047495 *
availability_30	1.94863	0.53128	3.668	0.000273 ***
calculated_host_listings_count	-0.71881	0.23872	-3.011	0.002748 **

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 97.93 on 457 degrees of freedom
Multiple R-squared:  0.4331,    Adjusted R-squared:  0.4244
F-statistic: 49.88 on 7 and 457 DF,  p-value: < 2.2e-16
```



Cleaning Data → All variables in model 1 → Eliminate multicollinearity by viewing the correlation table → Backward elimination → Selective variables in Model 2

2. Prediction: Multiple Regression Model


- i. Generate Equation
- ii. Predict Values
- iii. Values Match: 138 AUD



```
> predict(model2, myhouse)
1
138.1482
> -14.177 + 24.11*2 + 57.59*2 + 9.36*2 + -13.03*5 + 0.029*0 + 1.95*20 + -0.72*5
[1] 138.193
>
```

`visualize.t(stat=c(-1.987, 1.987), df=457,
section = "bounded")`

→ 0.952 of the shaded area →
very reliable



2. Prediction: Multiple Regression Model



R-Squared: 0.4363 → 0.4331

- sum of squared error divides by the sum of square total.
- Adjusted R-Squared: 0.4331 → 0.4244
- Less redundancy
- More efficient variations



RMSE

(the square root of mean squared error)

```
> accuracy(pred1, train$price)
              ME      RMSE      MAE      MPE      MAPE
Test set 1.404686e-12 97.0806 56.93222 -15.00545 35.50696
> pred2 <- predict(model2, valid)
> accuracy(pred2, valid$price)
              ME      RMSE      MAE      MPE      MAPE
Test set 5.709778 97.81584 57.33693 -10.39859 33.33211
> |
```

- Less error in train data: 97.0806 > 97.81584
- Not much differences in two RMSE

We have a RELIABLE model!




3. Classification: K-nearest Neighbors

Outcome: What kind of cancellation policy?

Variables:

- i. Host response rate
- ii. Price
- iii. Security Deposit
- iv. Review scores communication

##	k	accuracy
## 1	1	0.5018182
## 2	2	0.4472727
## 3	3	0.5309091
## 4	4	0.5454545
## 5	5	0.5672727
## 6	6	0.5563636
## 7	7	0.5781818
## 8	8	0.5927273
## 9	9	0.5818182
## 10	10	0.5781818
## 11	11	0.5745455
## 12	12	0.6145455
## 13	13	0.5745455
## 14	14	0.5818182



```
> # build knn model
> nn <- knn(train = train.norm.df[, 2:5], test = new.rental.norm.df,
  = 7)
> row.names(train.df)[attr(nn, "nn.index")]
[1] "109" "16" "394" "211" "273" "174" "133"
>
> new.rental.neighbor_valid <- train.df[c(109, 16, 394, 211, 273, 174, 133), ]
```

flexible (109), flexible (16), flexible (394), flexible (211), flexible (273), moderate (174), flexible (133)



FLEXIBLE

K-value: 12 | Accuracy: 61.45%
Flexible Cancellation Policy

```
> row.names(train.df)[attr(knn12, "nn.index")]
[1] "109" "16" "394" "211" "273" "174" "133" "161" "370" "223" "282" "106"
>
> new.rental.neighbor_valid <- train.df[c(109, 16, 394, 211, 174, 133, 161, 370, 223, 282, 106), ]
> new.rental.neighbor_valid$cancellation_policy #my policy is flexible
[1] flexible flexible flexible
[4] flexible moderate flexible
[7] flexible strict_14_with_grace_period moderate
[10] flexible moderate
```

3. Classification: Naïve Bayes

Outcome: What can measure to become instant bookable housing?

Variables:

- Host response rate
- Number of beds
- Available days in a month
- Cancellation Policy



```
fictional<-data.frame(beds=3,  
  cancellation_policy= s.factor("flexible"),  
  availability_30= 4,  
  host_response_rate=1)  
predict(naive,fictional)
```

```
[1] t  
Levels: f t
```

True for instant bookable!

Accuracy : 0.6419
95% CI : (0.6013, 0.681)
No Information Rate : 0.6211
P-Value [Acc > NIR] : 0.1621

Kappa : 0.1062

Accuracy : 0.6078
95% CI : (0.557, 0.6569)
No Information Rate : 0.5922
P-Value [Acc > NIR] : 0.285

Kappa : 0.0855

A-priori probabilities:
Y
f 0.3788927 t 0.6211073

Conditional probabilities:
cancellation_policy
Y flexible moderate strict strict_14_with_grace_period super_strict_30
f 0.237442922 0.333333333 0.000000000 0.406392694 0.000000000
t 0.105849582 0.242339833 0.000000000 0.649025070 0.000000000
cancellation_policy
Y super_strict_60
f 0.022831050
t 0.002785515

beds
Y [,1] [,2]
f 2.118721 1.460353
t 2.406685 1.517702
availability_30
Y [,1] [,2]
f 10.70776 9.918312
t 12.58217 8.624268
host_response_rate
Y [,1] [,2]
f 0.009538813 0.0013178383
t 0.009725905 0.0007377852

Training

validation

Results for the Conditions:

[accuracy 64.19% in training set]

- faster host response
- enough bed for at least one or two people
- more available housing,
- somehow flexible

3. Classification: Classification Tree



Purpose: Predict the size of the cleaning fee.





Variables:


- Room type
- Accommodates
- # of bathrooms, bedrooms and beds
- Price



**Bin the
Cleaning
Fees**



3. Classification: Classification Tree




Cross-Validation

- Determine the ideal size
- Find the minimum xerror
- Get the optimal cp value is 0.0039526

3. Classification: Classification Tree



Classification Tree with ideal size



For instance, the bottom-left-most terminal node
= “Medium” cleaning fee →
if it is the entire room or apartment and the
bedrooms < 2 and,
the price is not smaller than 120 and,
accommodates < 3 and,
the price is not smaller than 156.

3. Classification: Classification Tree

Confusion matrix

Training set

Confusion Matrix and Statistics

		Reference			
Prediction	No Fee	Low	Medium	High	
No Fee	0	0	0	0	
Low	14	72	18	3	
Medium	4	3	36	16	
High	26	6	74	231	

Overall Statistics

Accuracy : 0.674

95% CI : (0.6311, 0.7148)

No Information Rate : 0.497

P-Value [Acc > NIR] : 7.913e-16

Kappa : 0.4592

McNemar's Test P-Value : < 2.2e-16

Statistics by Class:

	Class: No Fee	Class: Low	Class: Medium	Class: High
Sensitivity	0.00000	0.8889	0.28125	0.9240
Specificity	1.00000	0.9171	0.93867	0.5810
Pos Pred Value	NaN	0.6729	0.61017	0.6855
Neg Pred Value	0.91252	0.9773	0.79279	0.8855
Prevalence	0.08748	0.1610	0.25447	0.4970
Detection Rate	0.00000	0.1431	0.07157	0.4592
Detection Prevalence	0.00000	0.2127	0.11730	0.6700
Balanced Accuracy	0.50000	0.9030	0.60996	0.7525

Validation set

Confusion Matrix and Statistics

		Reference			
Prediction	No Fee	Low	Medium	High	
No Fee	0	0	0	0	
Low	8	40	14	2	
Medium	6	3	13	10	
High	15	8	61	154	

Overall Statistics

Accuracy : 0.6198

95% CI : (0.5653, 0.6721)

No Information Rate : 0.497

P-Value [Acc > NIR] : 4.270e-06

Kappa : 0.357

McNemar's Test P-Value : 2.022e-14

Statistics by Class:

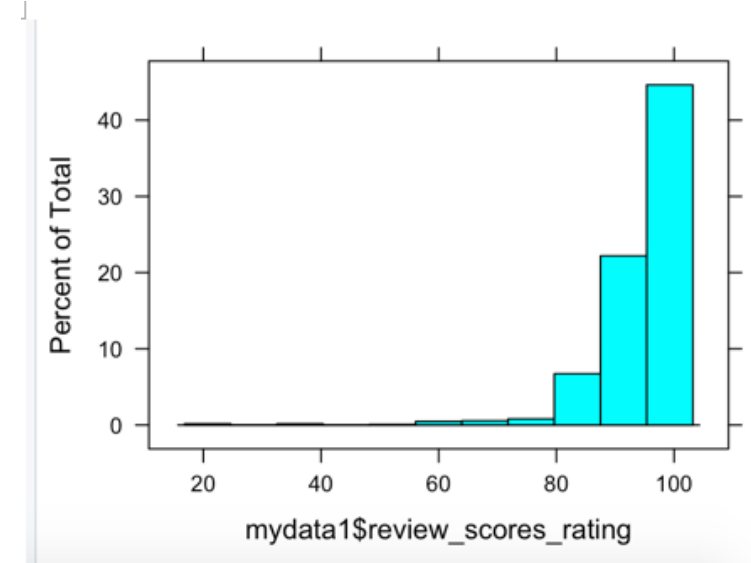

	Class: No Fee	Class: Low	Class: Medium	Class: High
Sensitivity	0.00000	0.7843	0.14773	0.9277
Specificity	1.00000	0.9152	0.92276	0.5000
Pos Pred Value	NaN	0.6250	0.40625	0.6471
Neg Pred Value	0.91317	0.9593	0.75166	0.8750
Prevalence	0.08683	0.1527	0.26347	0.4970
Detection Rate	0.00000	0.1198	0.03892	0.4611
Detection Prevalence	0.00000	0.1916	0.09581	0.7126
Balanced Accuracy	0.50000	0.8498	0.53525	0.7139



4. Clustering



- Special Data Cleaning: mean value → substitution → missing values
- Selected Variables: accommodates, bathrooms, bedrooms, beds, price, review score of rating, room type
- Feature Engineering:
 - Categorical Variables → Dummy Variables
 - Introduce “weight” to describe the related variables




```
> histogram(mydata1$review_scores_rating)
> summary(mydata1$review_scores_rating)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
20.00	92.00	96.00	94.23	99.00	100.00	302

```
> # a dataframe includes numerica variables
> accommodates <- mydata1$accommodates
> comfort <- mydata1$bathrooms*.35 + mydata1$bedrooms*.3 + mydata1$beds*.35
> price <- mydata1$price
> evaluation <- mydata1$review_scores_rating
> mydata2 <- data.frame(accommodates,comfort,price,evaluation)
```



4. Clustering



Optimal Cluster Groups: 5



4. Clustering



Categories	Attributes	Recommendations		
		Room Type	Surrounding	Promotions
Rich Couples	Enjoy high-level life	Private, Fancy, High reviewing score	Fine restaurant, Luxury places	Club member
Close Friends	Go out together	Entire room or Apartment, Big, High reviewing score	Nice place for taking pictures and hanging out	Discount for restaurants, Uber or Lift
Thrifty People	Save Money	Entire room or Apartment, Lower price	Free places	Awards from completing task
Colleagues	Go out for business	Shared	Convenient Transportation	Discount for room service
Family	Go out together	Entire room or Apartment, Big, High reviewing score	Nice place for taking pictures and hanging out	Discount for restaurants, renting car

Conclusion & Recommendations



- Value Privacy
- Suitable for Vacations
- Advice for Host: Flexible Cancellation, Faster Host Response, Cheaper Housing
- Relaxing Neighborhood → Investment Opportunities
- More Diverse Property Type



Any Questions?